

A group of yellow Minions with large goggles are gathered in a modern office with large windows and industrial-style lighting. They are looking up at a glowing, swirling sphere of blue and yellow light that appears to be made of energy or data. One Minion in the center foreground is holding a sign that says "AI".

A PRACTICAL AI KEYNOTE

Not Everything Is AI

A 60-minute field guide to what's real, what's hype, and when you actually need a machine that "thinks".

"Draw an author on a plane using the wifi"

Ethan Mollick ran the *same* prompt through the best AI image model, year after year. Watch how fast the ground moved.

□ 2022
drop image


□ 2023
drop image

□ 2024
drop image

□ 2025
drop image

The tension: dizzyingly fast — *and* drowning in hype. Today we separate the two.



A group of yellow Minions in blue overalls are in a classroom. One Minion is standing on a desk on the left, looking at a whiteboard. Another Minion is on the right, also looking at the whiteboard. Several other Minions are standing in the foreground, looking up at the whiteboard. The whiteboard has a diagram titled 'MY ROADMAP' with three steps: 'STEP 1: IDEATE', 'STEP 2: GATHER BANANAS', and 'STEP 3: BUILD!'. Arrows connect the steps in a sequence. There are also some smaller boxes with text like 'STEP: S!' and 'STEP: S!' on the whiteboard. The background shows a classroom with posters on the wall and a window on the right.


What we're doing in this hour

A practical tour of what "AI" actually is. The goal isn't to put AI on everything — it's to point the *right tool* at the *right problem*.

The plan

- **The honest tour** — from basic Python up to ChatGPT, and why "AI" is the top of a tall ladder
- **Kahoot** — see what stuck (and have some fun)
- **Live demo** — a real pipeline where the LLM is the *last 5%*, not the whole thing

Team meeting, not a lecture. Interrupt me.



Is it all AI? Is it all hype?

Neither. The honest answer is a **ladder**. Most of the value in our work sits on the *lower rungs* — the boring, reliable, auditable ones.

The mistake everyone is making right now: reaching for the top rung (an agent, an LLM) for a problem a 10-line script already solved.

All-or-nothing is a fallacy. Let's climb the ladder so you can tell the rungs apart.



The ladder — five rungs

1. **Classic automation** — "if this, do that" (basic Python)
2. **NLP** — find patterns in text
3. **Machine Learning** — learn rules from examples
4. **The Transformer** — the engine under modern AI
5. **ChatGPT** — what happens when you feed it the internet

One image, one idea per rung. You've already used most of these without knowing the name.



Rung 1 — Classic automation

A computer doing **exactly what you told it**, fast, forever, without complaining. No "intelligence" at all.

- Rename 1,000 files · merge 12 Excels · send the email when the report lands
- 100% predictable · 100% auditable · costs nothing to run

This solves more of your day than any LLM will. It just doesn't trend on LinkedIn.

ANCIENT MAP

ENIGMA

UNDERSTAND

CLUES

Rung 2 — NLP (language as data)

Teaching a computer to **find structure in text** — without it "understanding" anything.

- Pull the total out of 500 invoices · flag emails that mention "complaint"
- Classify a document as *contract* vs *invoice* vs *other*

Still mostly rules and patterns. Reliable, cheap, explainable.

ANCIENT



Rung 3 — Machine Learning

Instead of writing the rules, you **show it examples** and it learns the rule itself.

- "Here are 10,000 customers who churned" → predict who's next
- Propensity, scoring, fraud, forecasting — the *workhorse* of analytics

Older than you'd think (the idea is from the 1950s). **Most "AI projects" are really this** — and that's good: it's measurable and auditable.

A Minion character is sitting on a stool in a server room, surrounded by glowing orange cables. Above the Minion, several words are displayed in glowing orange neon-like font, connected by glowing orange lines. The words are: "cat", "sat", "mat", "on", "the", "the", "fox", "runs", "mat".

cat sat mat on the the fox runs mat

Rung 4 — The Transformer (2017)


One architecture change — "**attention**" — let models weigh which words matter to which. That's the engine under every modern LLM.

You don't need the maths. You need one fact: this is the part that made the next rung possible — and it's only *eight years old*.

So how is ChatGPT built? (Karpathy's arc)

Step 1 — Pre-training: read essentially *all the public text on the internet* and learn to predict the next word. You get an **"internet simulator"** — brilliant, but it just autocompletes anything.


□ Optional: diagram — "internet → next-word predictor"
drop image / keep gradient

A Minion character is dressed as a hotel concierge, wearing a purple uniform with gold buttons and a matching cap. The Minion is standing behind a dark wooden reception desk. On the desk, there is a brass bell on the left and a brown leather-bound book titled "Guest Book" on the right. The background is a grand hotel lobby with wood-paneled walls, large windows with white curtains, and several wall-mounted lamps. A chandelier hangs from the ceiling. The Minion's hands are clasped in front of it.

Step 2 — Post-training: make it an assistant

Show it thousands of examples of *helpful, well-behaved answers*. Now the internet-simulator stops rambling and starts **acting like an assistant**.

It hasn't learned new facts — it's been **biased toward being useful**.




Step 3 — RLHF: turn it into a specialist

Humans rank answers good vs bad, and the model is rewarded for the good ones. [Reinforcement Learning from Human Feedback](#) sharpens the assistant into a specialist.

read the internet → internet simulator → assistant → specialist. That's the whole arc.

▶ Karpathy — How ChatGPT works (backup clip)



So — when do you actually need the top rung?

Rarely. The LLM is a brilliant, expensive, occasionally-hallucinating generalist. Use it where nothing else works: open-ended language, judgement, messy edge cases.

- Need a deterministic, auditable answer? → **lower rung**
- Need to handle the unpredictable 5%? → **top rung, as a scalpel**

Hold that thought — the demo is exactly this.

ENERGY RESET

Kahoot time

10 questions. Some real, some ridiculous. Phones out.

▶ Launch Kahoot

Import assets/kahoot.csv into the Kahoot template first.



The demo: automated document triage

A real pipeline that reads incoming PDFs, classifies them, applies a set of rules, and *only at the very end* calls an LLM — as a judge, to catch false positives.

Watch which rung does the heavy lifting.

One problem, four rungs — in order



The LLM is the **last 5%** — not the whole pipeline.


□ Optional: live mockup screenshot / screen-share goes here



"I want AI to do my laundry and dishes so that I can do my art and writing — not for AI to do my art and writing so that I can do my laundry and dishes."

— Joanna Maciejewska

Point the machine at the boring, valuable work. Keep the judgement for us.

A group of Minions stands on a rocky cliff at sunset. A large white flag is planted in the ground, and the sun is low on the horizon, casting long shadows. The scene is framed by a dark semi-transparent box containing text.

The takeaway

Pick the right rung. Keep the boring automation un-hyped and reliable, reserve AI & agents for where they genuinely earn their cost, and keep all of it auditable.

Right tool, right problem. Every time.

Questions? Now's the time.

